

Problem Set #1

MACS 30250, Dr. Evans

Due Monday, May. 4 at 1:30pm

1. **1D kernel density estimator (5 points).** The [COVIDincubation.txt](#) file is a comma-delimited data file that includes 59 observations on individuals in China who tested positive for the COVID-19 virus. The data contain the following three variables: `gender` (“M” or “F”), `age` (integers between 10 and 70), and `symp_days` (days until symptomatic, float). This is a subset of the variables in the dataset used by [Men et al. \(2020\)](#). The `symp_days` variable represents the incubation period for each individual, or the number of days until symptoms were manifest.
 - (a) Create three histograms, each of `symp_days` (Incubation period, days to symptomatic). The first one is the overall histogram. Let each histogram have 15 bins over the range of days from 2 to 15 (the maximum in the data). In the `matplotlib.pyplot.hist()` call, set the histogram to density `density=True`. Let the first histogram be for all the data. Let the second histogram be for individuals of `age ≤ 40`, and let the third histogram be for individuals `age > 40`.
 - (b) Fit a Gaussian KDE as an approximation of the incubation period for each of the three subsets of data from part (a). Use the `GridSearchCV` and `LeaveOneOut` methods as in the [VanderPlas notebook](#) to choose an optimal bandwidth, and report your optimal bandwidths for the three KDEs. For your grid search, use 500 exponentially spaced bandwidths between 0.1 and 10 using the code: `bandwidths = 10 ** np.linspace(-1, 1, 500)`. Plot each of the KDE distributions in one plot with a legend that shows which is which. This figure should look like the right panel of Figure 2 in [Men et al. \(2020\)](#).
 - (c) What does this tell you about COVID-19 incubation periods of young versus old individuals?
2. **2D kernel density estimator (5 points).** This exercise uses two data files: [BQ_ind_data.txt](#) and [BQ_probmat.txt](#). The first data file ([BQ_ind_data.txt](#)) contains 70,000 observations on two variables: `age` and `income_pct1`. These represent individual data where each observation is an individual that represents a person who received an equal bequest. So age and income groups that received more bequests will be represented by more individuals (observations). The second data file ([BQ_probmat.txt](#)) contains the empirical histogram information. It is a 73×7 matrix of percentages representing the values of a two-dimensional histogram of the percent of the U.S. population that receives all the bequests (inheritances) by a recipient’s age (ages 18 to 90, rows) and by

a recipient's lifetime income group (7 categories, columns). The seven lifetime income groups are percentiles. Let $prcntl_j$ be the percent of the population in lifetime income group j . The lifetime income groups in the $J = 7$ columns of the `BQ_probmat.txt` data are the following.

$$prcntl = [0.25, 0.25, 0.20, 0.10, 0.10, 0.09, 0.01], \quad \text{such that} \quad \sum_{j=1}^7 prcntl_j = 1$$

You can read this file into memory using the `numpy.loadtxt` function.

```
bq_data = np.loadtxt('BQ_probmat.txt', delimiter=',')
```

So the `[11, 5]`-th element of the `bq_data` matrix represents the percent of total bequests (inheritances) received by age-28 and lifetime income group $j = 5$ (80th to 90th percentile of lifetime income).

- (a) Read in the `BQ_probmat.txt` data as a 73×7 NumPy array. Plot the 2D empirical histogram of these data as a 3D surface plot with age and income group on the x -axis and y -axis and the histogram density on the z -axis using a 3D surface plot tool (not a 3D bar histogram tool). Make sure that the axes are labeled correctly. And make sure that your 3D histogram is presented from a perspective that allows a viewer to see that data (don't let the data be hidden by a poor angle of the plot.)
- (b) Use the `BQ_ind_data.txt` data to fit a bivariate Gaussian kernel density estimator to the data using the `scipy.stats.gaussian_kde` method. Choose a bandwidth parameter λ that you think is best. Justify your choice of that parameter. Your justification should have to do with the tradeoff between overfitting (too low a value) and underfitting (too high a value). Plot the surface of your chosen kernel density estimator. Make sure that the axes are labeled correctly. And make sure that your 3D histogram is presented from a perspective that allows a viewer to see that data. What is the estimated density for bequest recipients who are age 61 in the 6th lifetime income category ($j = 6$, 90th to 99th percentile).

References

Men, Ke, Xia Wang, Yihao Li, Guanwei Zhang, JingjingHu, Yanyan Gao, and Henry Han, "Estimate the Incubation Period of Coronavirus 2019 (COVID-19)," February 2020. Unpublished.